Exploratory Data Analysis for Business Intelligence

961731 Business Data Analytics

Asst. Prof. Dr. Prompong Sugunnasil

13/07/2026

Module 1: Enterprise Data Understanding & CRISP-DM

Aligning technical data discovery with strategic corporate objectives.

What is CRISP-DM?

- CRISP-DM (Cross-Industry Standard Process for Data Mining) is the global standard framework for structuring analytical projects.
- Rather than a linear academic task, it is a cyclical business process that starts and ends with business value.
- Data Understanding is the phase where we audit, clean, and explore raw data to verify its quality before building predictive models.

Stage	Focus	Key Output
1. Business Understanding	Goals & Objectives	Project Plan
2. Data Understanding	Ingest & Explore	Data Quality Audit
3. Data Preparation	Clean & Format	Clean Dataset
4. Modeling	Algorithm Selection	Predictive Model
5. Evaluation	Business Validation	ROI Decision
6. Deployment	Integration	Operational Dashboard

Align CRISP-DM to Strategy

- Business objectives must define the parameters of data exploration to prevent teams from mining irrelevant correlations.
- Data scientists must map technical variables directly to income statement lines and balance sheet assets.
- This alignment avoids analysis paralysis and ensures that downstream models produce measurable EBITDA gains.

Corporate Objective	Operational Lever	Target Data Source
Increase Profitability	Reduce Customer Churn	Contract Billing Records
Free Up Cash Flow	Minimize Lead Time Volatility	Warehouse Dispatch Logs
Reduce Acquisition Cost	Optimize Marketing Target	Ad Campaign Response Data

CHIANG MAI UNIVERSITY

Operationalize Data Understanding

- Running a systematic data audit exposes structural defects, system recording delays, and missing customer records.
- Evaluate raw data across four dimensions: completeness, accuracy, consistency, and timeliness.
- Early detection of quality issues reduces downstream modeling rework by 40% and protects the business from making bad decisions.

Dimension	Focus	Risk if Ignored
Completeness	Missing values and rows	Biased sample populations
Accuracy	Correctness of records	Incorrect model outputs
Consistency	Same format across sources	Joins fail or corrupt data
Timeliness	Latency of data entry	Models using outdated facts

CHIANG MAI UNIVERSITY

Drive Hypothesis-Led Discovery

- Hypothesis-led exploration focuses scarce analytical resources on verifying specific, high-potential business opportunities.
- Construct MECE (Mutually Exclusive, Collectively Exhaustive) issue trees to break a large business KPI into testable components.
- This structured process increases analyst productivity by 50% by avoiding open-ended, aimless data mining.

KPI Level	Driver Branch	Diagnostic Question
Profit Decline	Revenue branch	Is price decreasing, or is volume declining?
Revenue	Volume branch	Which customer segments have stopped purchasing?
Cost	Variable branch	Has shipping volatility increased product costs?

Map Operational Levers

- Value-driver trees mathematically connect low-level operational variables to high-level financial outcomes (e.g., ROIC or Net Income).
- Deconstruct profit into its component revenue drivers (prices, volumes) and cost drivers (fixed, variable, overheads).
- Interactive value-driver models enable real-time 'what-if' simulations to prioritize data science projects based on financial impact.

Financial KPI	Operational Driver	Data Science Lever
EBITDA	Order Frequency	Churn Prediction Model
EBITDA	Shipping Cost	Route Optimization Model
EBITDA	Average Order Value	Product Recommender System

Supply Chain & Working Capital

- Dataset Link: https://www.kaggle.com/datasets/olistbr/brazilian-ecommerce
- Business Context: CFO mandate to reduce delivery lead times to free up working capital.
- Analytical Task:
  - Deconstruct order delivery times into shipping, warehousing, and transit phases.
  - Quantify the relationship between shipping delays and customer repeat purchase rate.
  - Translate a 2-day reduction in average transit time into working capital savings and EBITDA impact.

Module 2: Descriptive Analytics & Baseline Operations

Establishing baseline performance and measuring process volatility.

What are Averages and Spread?

- Averages (Mean, Median, Mode) try to find the 'typical' value in a mountain of data.
- Spread (Range, Variance, Standard Deviation) measures how much the individual values scatter or vary around that typical center.
- For business decisions, knowing the average is not enough; we must understand the variance to manage operational risk.

Metric Group	Core Metrics	Business Analogy
Averages	Mean (average), Median (midpoint), Mode (most common)	What is our standard customer transaction size?
Spread	Range (min-max), Std Dev (volatility), IQR (middle 50%)	How volatile is our transaction size from day to day?

Establish Baselines via Averages

- Central tendency metrics establish the baseline for standard transactions, customer lifetime, and process cycles.
- Comparing the mean and median reveals the skewness of the distribution and the presence of extreme customer segments.
- Using the median instead of the mean prevents strategic pricing and inventory errors when dealing with highly skewed spend data.

Distribution Type	Mean vs. Median	Business Context
Symmetric	Mean = Median	Uniform transaction sizes
Positively Skewed	Mean > Median	Skewed by high-value outliers (e.g. e-commerce spend)
Negatively Skewed	Mean < Median	Skewed by low-value outliers (e.g. discounted sales)

Quantify Risk via Volatility

- Spread metrics quantify operational volatility, serving as a direct measure of process stability and service reliability.
- Calculate standard deviation and interquartile range (IQR) to identify the consistency of manufacturing or shipping cycles.
- Tightening the spread allows operations to meet Service Level Agreements (SLAs) and avoid contract penalties.

Metric	Mathematical Concept	Business Application
Range	Difference between Max and Min	Simplest measure of process extremes
Std Dev	Average distance from the mean	Benchmark for process quality (e.g. Six Sigma)
IQR	Spread of the middle 50% of data	Volatility baseline ignoring outliers

Monetize Volatility & Quality

- Volatility in supplier lead times forces companies to carry expensive safety stock, tying up cash on the balance sheet.
- Use the standard deviation of lead time to mathematically calculate safety stock levels required for a 95% service level.
- Reducing lead-time volatility directly reduces holding costs, frees up working capital, and improves cash flow.

Lead Time Volatility	Required Safety Stock	Annual Holding Cost
High (5 Days)	1,200 Units	$120,000
Medium (2 Days)	500 Units	$50,000
Low (0.5 Days)	125 Units	$12,500

Retail Baselines & Margin Defense

- Dataset Link: https://www.kaggle.com/datasets/olistbr/brazilian-ecommerce
- Business Context: CEO wants to standardize regional sales performance and optimize pricing strategies.
- Analytical Task:
  - Calculate central tendency and spread (mean, median, IQR) of order values by product category and region.
  - Identify high-volatility categories and calculate the margin risk.
  - Formulate a pricing recommendation that reduces revenue variance and increases EBITDA by 3%.

Module 3: Visual Diagnostic Analytics & Summary Stat Failures

Leveraging visualization to expose the critical limits of summary statistics.

Why Do We Visualize Data?

- The human brain processes visual information 60,000 times faster than text or tables of numbers.
- Summary statistics (like average and correlation) compress complex data into single numbers, but they discard critical shape information.
- Visualization acts as a diagnostic lens that reveals structural patterns, anomalies, and relationships that statistical formulas miss.

Representation	Human Cognitive Load	Pattern Recognition Speed
Table of Numbers	Extremely High (requires comparison)	Very Slow (> 10 seconds)
Visual Chart	Very Low (exploits visual cortex)	Instant (< 250 milliseconds)

X1		X2		X3		X4	
x	y	x	y	x	y	x	y
10.00	8.04	10.00	9.14	10.00	7.46	8.00	6.58
8.00	6.95	8.00	8.14	8.00	6.77	8.00	5.76
13.00	7.58	13.00	8.74	13.00	12.74	8.00	7.71
9.00	8.81	9.00	8.77	9.00	7.11	8.00	8.84
11.00	8.33	11.00	9.26	11.00	7.81	8.00	8.47
14.00	9.96	14.00	8.10	14.00	8.84	8.00	7.04
6.00	7.24	6.00	6.13	6.00	6.08	8.00	5.25
4.00	4.26	4.00	3.10	4.00	5.39	19.00	12.50
12.00	10.84	12.00	9.13	12.00	8.15	8.00	5.56
7.00	4.82	7.00	7.26	7.00	6.42	8.00	7.91
5.00	5.68	5.00	4.74	5.00	5.73	8.00	6.89

Look Beyond Averages

- Relying solely on summary statistics can lead to strategic failures when completely different datasets yield identical averages.
- Technical teams must always plot raw data to confirm that the statistical models match the actual operational reality.
- Expose hidden distributions (such as linear, curved, or clustered patterns) that share identical means, variances, and correlations.

Property	Value
Mean of x	9
Sample variance of x	11
Mean of y	7.50
Sample variance of y	4.125
Correlation between x and y	0.816
Linear regression line	y = 3.00 + 0.500x

Visualize Raw Data

- Automated reporting pipelines often suffer from confirmation bias by failing to flag data that conforms to expected statistical averages.
- Complex datasets, like the Datasaurus Dozen, show that data can take wild geometric shapes while maintaining identical summary statistics.
- Data visualization acts as an operational defense mechanism to detect system bugs, data shifts, and new customer segments.

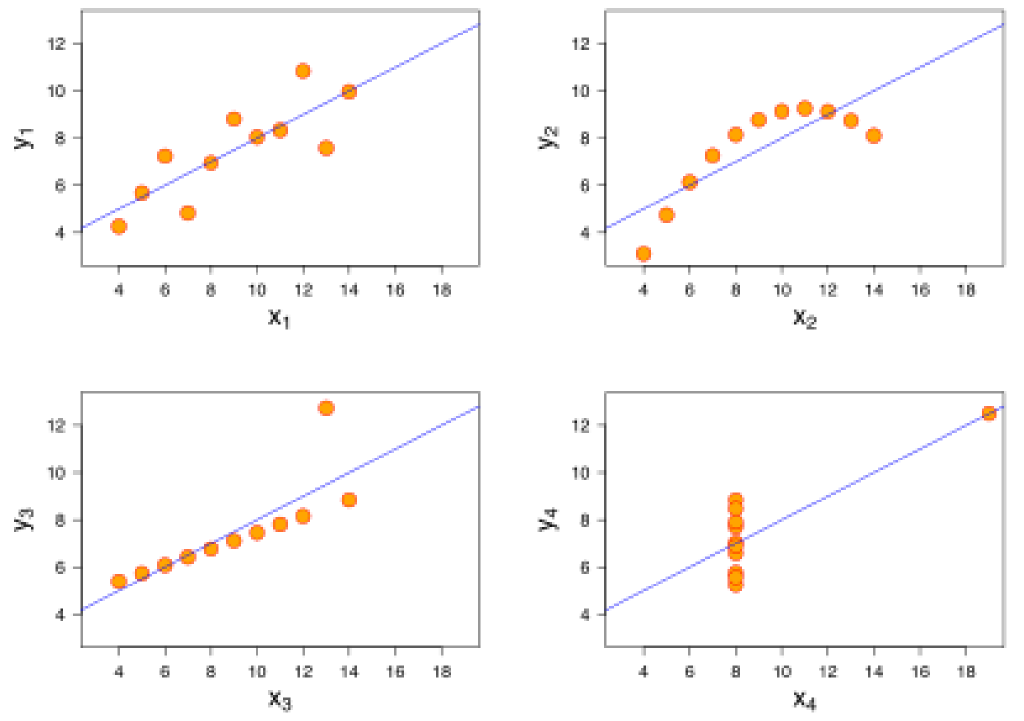

Identify Anomalies via Histograms

- Histograms group continuous data into ranges (bins) to reveal the shape and frequency of process variables.
- Exposing bimodal distributions reveals that a single process contains two distinct operating modes or customer segments.
- Adjusting bin size is critical: too few bins hide operational patterns, while too many bins create distracting noise.

Distribution Shape	Operational Indication	Strategy Action
Unimodal (Normal)	Stable, standard process	Optimize parameters
Bimodal (Two Peaks)	Two different systems mixed	Segment and analyze separately
Skewed (Long Tail)	Outliers and system delays	Implement process controls

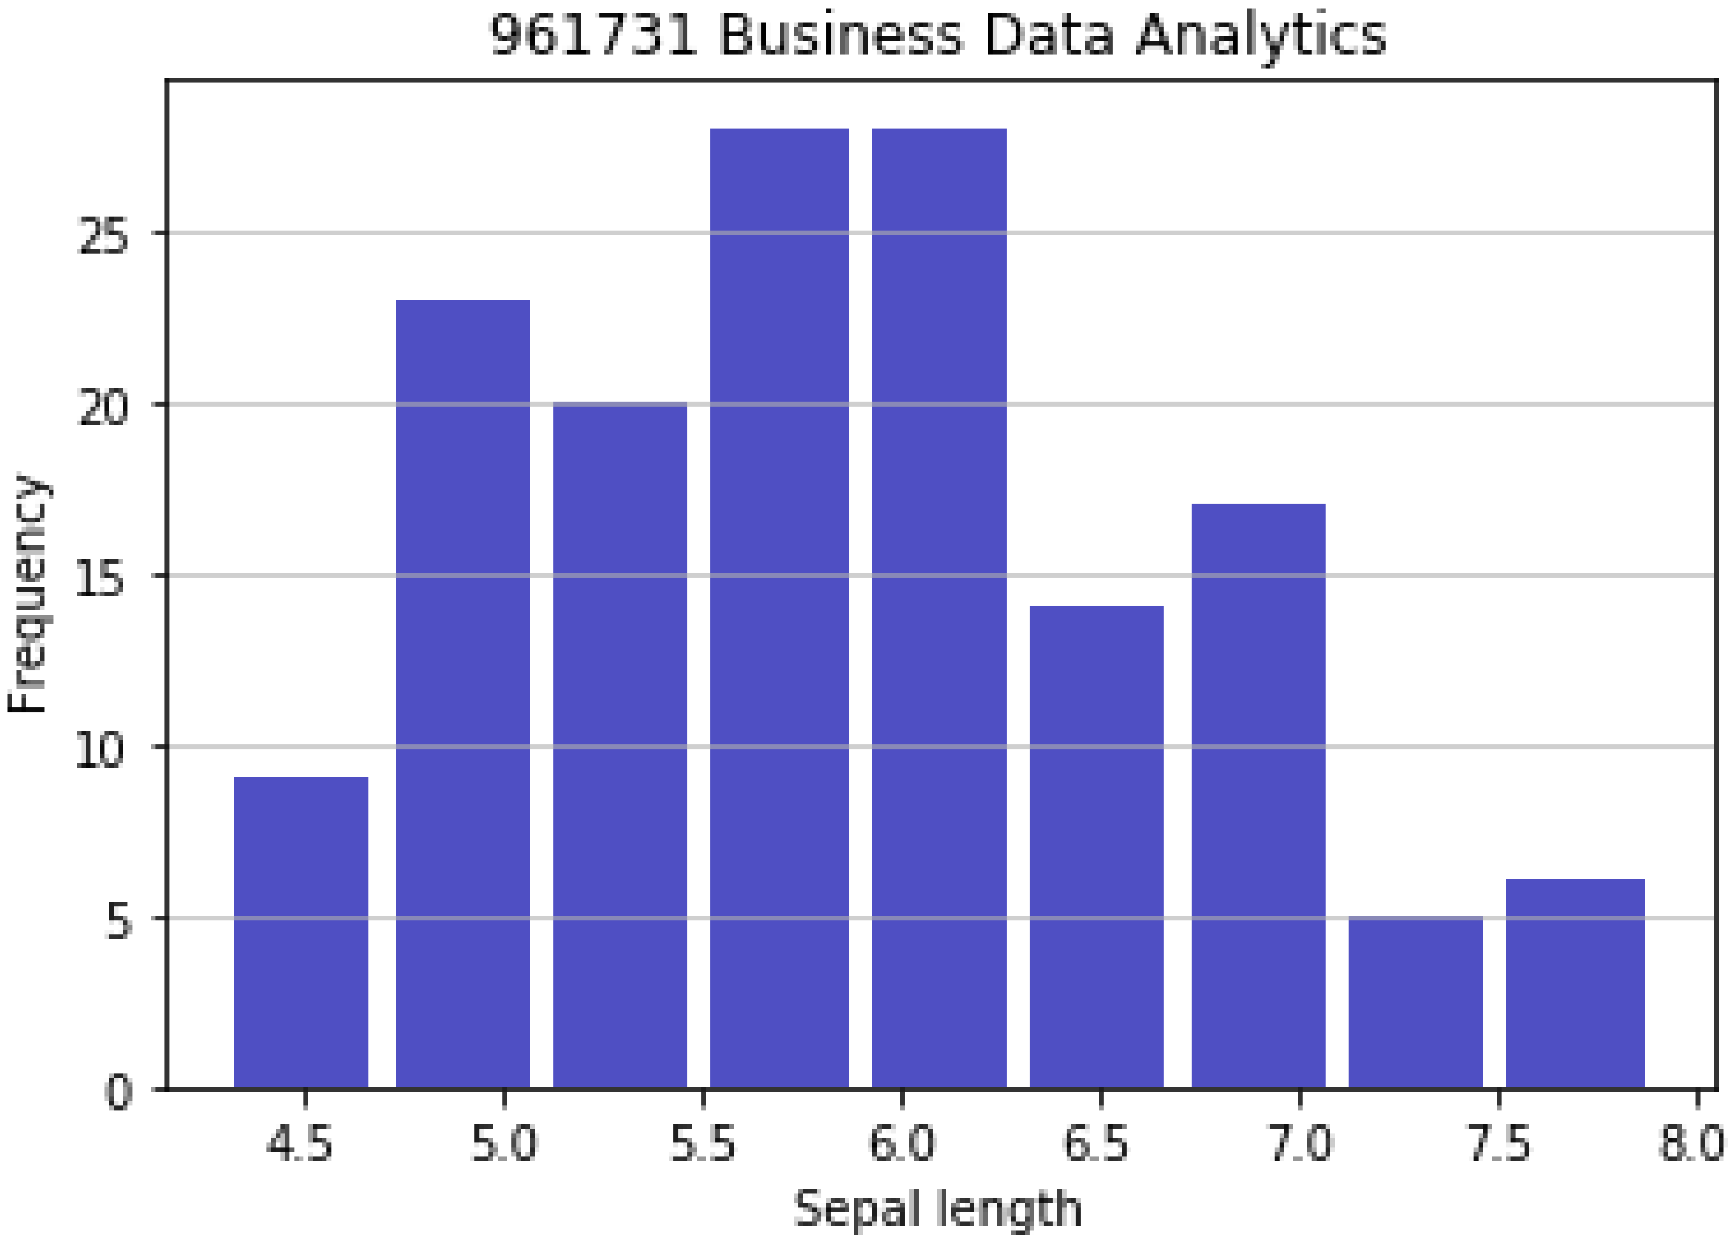
961731 Business Data Analytics
Frequency
Sepal length
0
5
10
15
20
25
4.5
5.0
5.5
6.0
6.5
7.0
7.5
8.0

CHIANG MAI UNIVERSITY

Control Variance via Boxplots

- Boxplots visually summarize a distribution using quartiles, showing the median, the middle 50% (IQR), and extreme outliers.
- Boxplots are the premier tool for comparing performance across multiple categories, stores, or regional hubs.
- Outliers on a boxplot highlight process failures, system errors, or service breaches that violate SLAs.

Visual Element	Statistical Definition	Operational Meaning
Center Line	Median (Q2)	The center of process speed
Box	Interquartile Range (Q1 to Q3)	Normal range of operations
Whiskers	1.5 * IQR from Box	Acceptable process boundary
Isolated Points	Outliers	Service breaches & exceptions

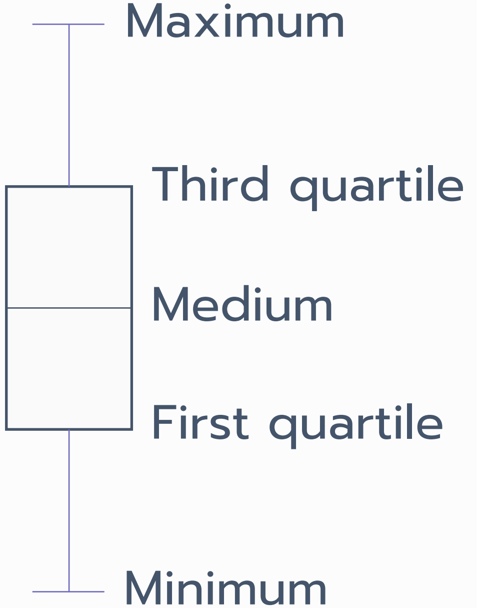

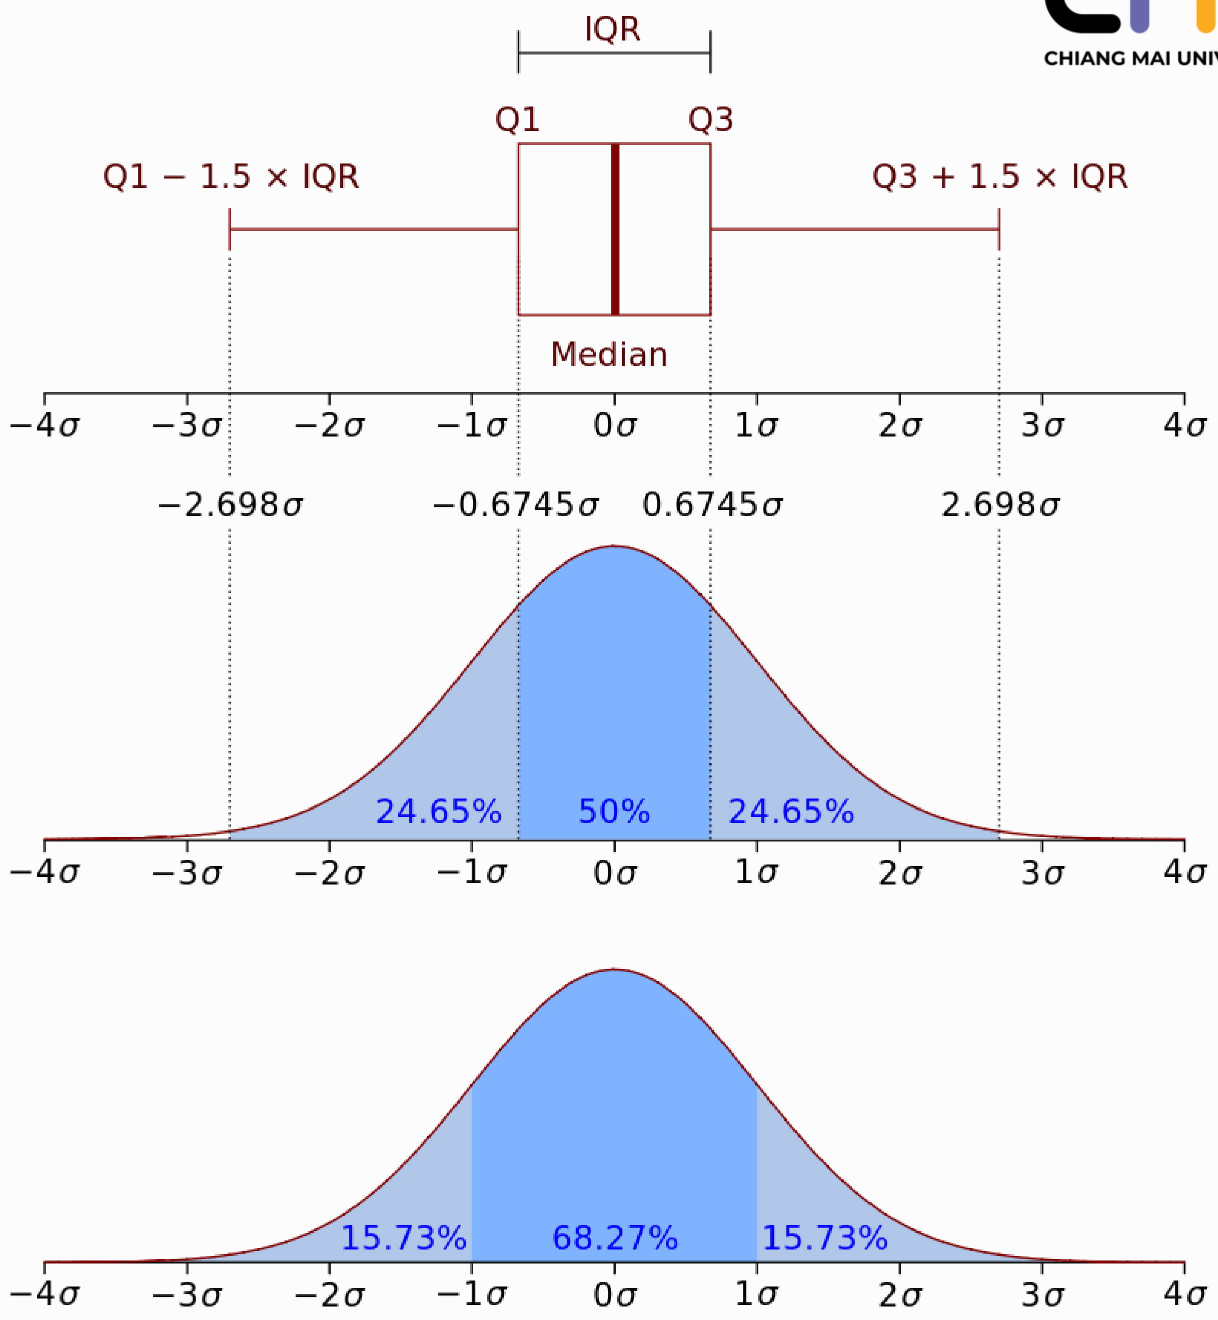

https://en.wikipedia.org/wiki/Interquartile_range

Map Linear & Non-Linear Relations

- Scatterplots plot individual data points along two axes to reveal the strength and direction of relationships between continuous variables.
- Expose non-linear relationships, such as the diminishing returns of marketing spend or saturation levels in operational throughput.
- Analysts must use scatterplots to screen for outliers that heavily distort linear correlation calculations.

Visual Pattern	Correlation Coefficient	Business Interpretation
Downward Linear	Near -1.0	Strong inverse driver (e.g. price vs. volume)
Curved (Logarithmic)	Misleadingly low linear r	Diminishing returns (marketing saturation)
Flat / Random	Near 0.0	No relationship; stop investing

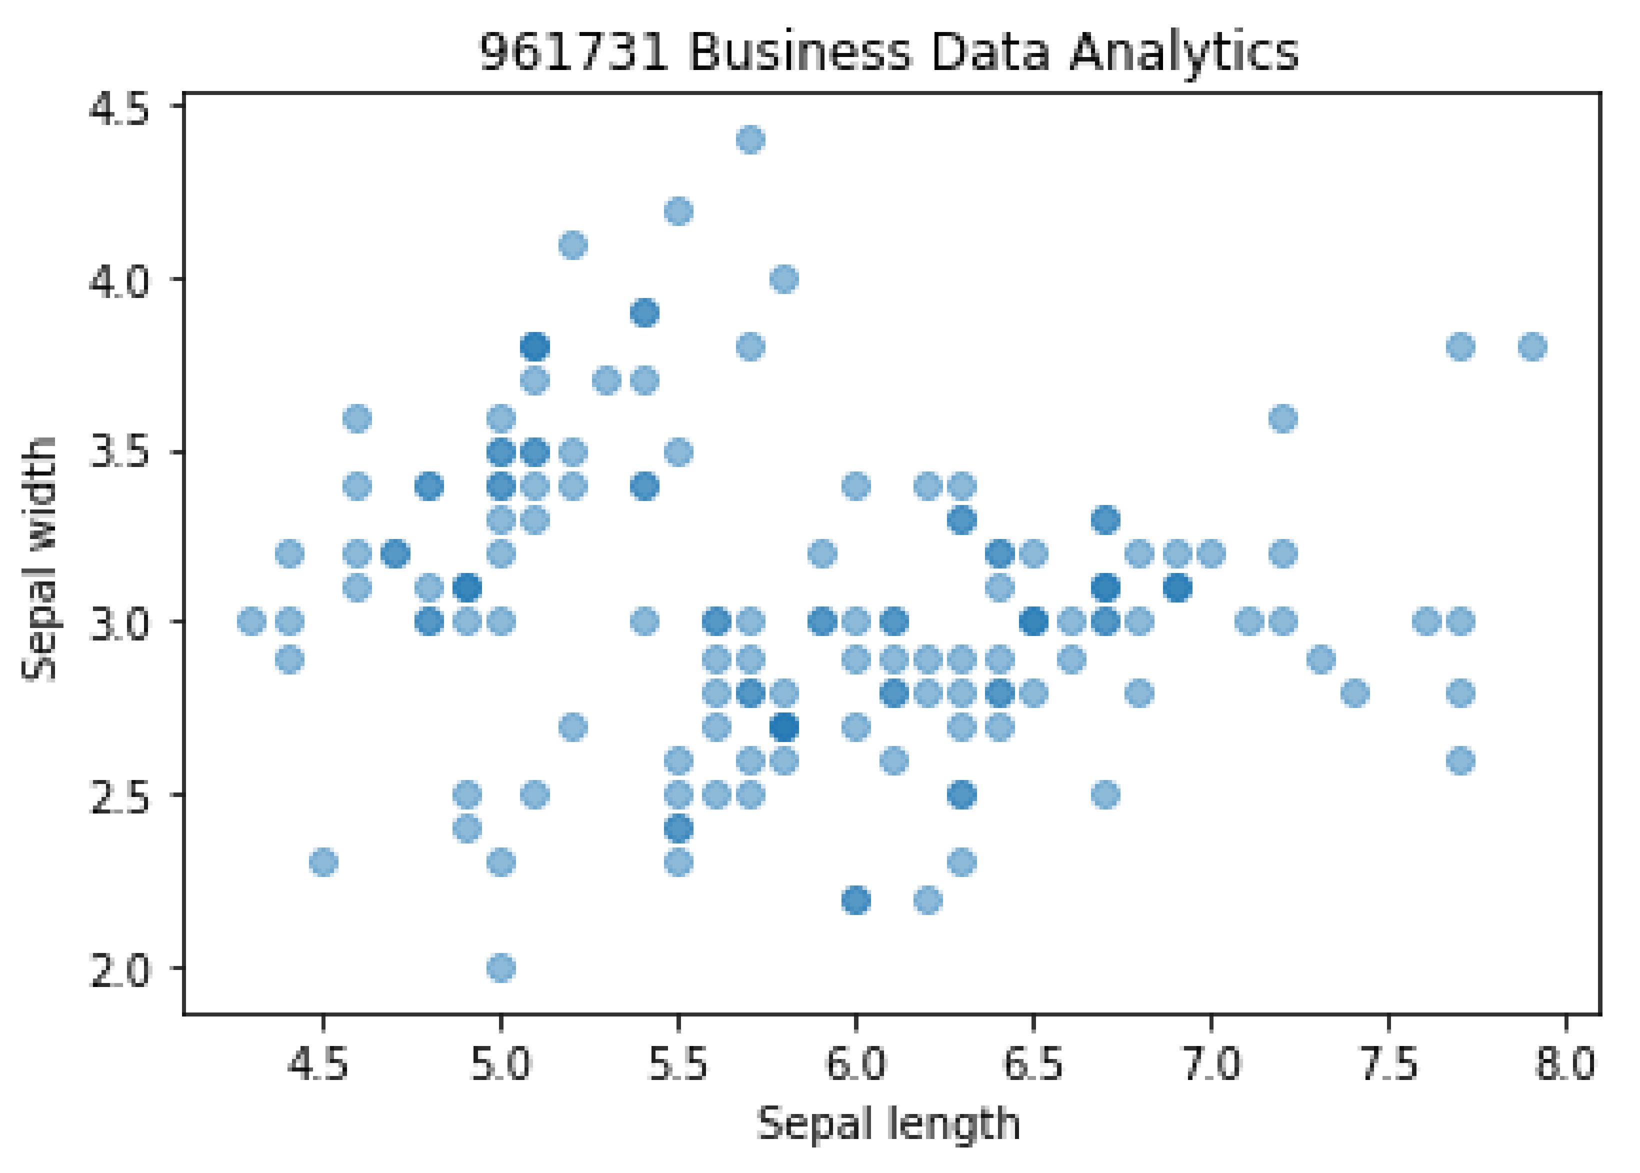
961731 Business Data Analytics
Sepal width
Sepal length
4.5
4.0
3.5
3.0
2.5
2.0
4.5
5.0
5.5
6.0
6.5
7.0
7.5
8.0

SLA Diagnostics & Penalties

- Dataset Link: https://www.kaggle.com/datasets/stephaniekerr/datasaurus-dozen
- Business Context: Server latency SLAs are tied to high penalties; average latency looks compliant, but customer complaints are spiking.
- Analytical Task:
  - Plot the distributions of system latency using histograms and boxplots.
  - Identify the hidden bimodal distribution and outliers that represent SLA breaches.
  - Calculate the financial penalty of these breaches and compare it to the cost of system upgrades.

Module 4: Categorical Diagnostics & Operational Insights

Uncovering structural business drivers through cross-tabulation and significance tests.

What is Cross-Tabulation?

- Categorical variables represent groups or labels (e.g., Country, Contract Type, Churn Status: Yes/No) rather than numbers.
- Cross-tabulation (also called a contingency table) is a grid that counts how often combinations of categories occur together.
- It allows us to see if one category is strongly associated with another (e.g., do paperless billing customers churn more?).

Customer Segment	Renewed	Cancelled	Cancel Rate
Month-to-Month	2,220	1,655	42.7%
One-Year Contract	1,307	166	11.3%
Two-Year Contract	1,647	48	2.8%

CHIANG MAI UNIVERSITY

Uncover Structural Drivers

- Cross-tabulation reveals how categorical variables interact, serving as the foundation for customer segmentation and risk profiling.
- Joint frequencies and percentages expose disproportionate churn or purchase patterns in specific product tiers.
- Identifying these high-risk categorical combinations enables targeted marketing, optimizing Customer Acquisition Cost (CAC).

Category 1 (Contract)	Category 2 (Billing)	Joint Churn Risk	Retention Priority
Month-to-Month	Paperless	Critical (45%)	High
Month-to-Month	Mailed	High (32%)	Medium
Two-Year Contract	Mailed	Very Low (1.5%)	Low (Self-sustaining)

Validate via Chi-Square Tests

- A Chi-Square test determines whether an observed relationship between categorical variables is statistically significant or merely random noise.
- The test compares observed cell counts against the expected counts under the assumption of complete independence.
- Using significance tests prevents organizations from launching expensive marketing campaigns based on coincidental, non-replicable patterns.

p-value	Significance Level	Business Decision Action
< 0.05	Statistically Significant	Segment behaves differently; launch targeted campaign
>= 0.05	Not Significant	Observed difference is random noise; do not modify strategy

Screen Features via Heatmaps

- Correlation heatmaps visually filter dozens of variables to quickly find the primary drivers of business outcomes.
- A diverging color scale highlights strong positive, strong negative, and near-zero linear correlations.
- Screening the feature space helps identify multi-collinearity (independent variables that are highly correlated with each other), simplifying models.

Value Range	Correlation Strength	Modeling Guideline
+0.7 to +1.0	Strong Positive	High predictive value, but check for redundancy
-0.3 to +0.3	Weak or None	Low individual linear predictability
-0.7 to -1.0	Strong Negative	High predictive value (inverse driver)

Telco Retention & CLV

- Dataset Link: https://www.kaggle.com/datasets/blastchar/telco-customer-churn
- Business Context: CMO wants to reduce customer churn from 26% to under 15% using targeted contract and payment incentives.
- Analytical Task:
  - Construct cross-tabulations of Churn against Contract, PaymentMethod, and PaperlessBilling.
  - Conduct Chi-Square tests to validate statistical significance.
  - Map the results to a financial impact matrix showing the EBITDA uplift of converting month-to-month contracts to 1-year contracts.

Module 5: Scope & Depth of Exploratory Analysis

Navigating the levels of data exploration from single-variable baselines to complex multivariate diagnostics.

CHIANG MAI UNIVERSITY

Hierarchy of Exploratory Depth

- Exploratory depth ranges from simple descriptive summaries of a single variable to multi-variable interaction analysis.
- Level 1 (Univariate) answers 'What does one variable look like?';
- Level 2 (Bivariate) answers 'How do two variables relate?'.
- Level 3 (Multivariate) and Level 4 (Hypothesis-Driven) isolate the true drivers of business outcomes by controlling for confounding factors.

Exploratory Level	Diagnostic Focus	Core Business Question
Level 1: Univariate	Distribution of a single variable	What is the typical customer age or order value?
Level 2: Bivariate	Relationship between two variables	Do paperless-billing customers churn more often?
Level 3: Multivariate	Interaction of three or more variables	How does monthly spend impact churn across contract tiers?
Level 4: Hypothesis-Driven	Root cause testing using business priors	Is our profit margin decline driven by shipping costs?

Univariate & Bivariate Scope

- Univariate diagnostics establish operational boundaries, scanning for data entry errors, null values, and skewness.
- Bivariate diagnostics map simple associations using correlation matrices, scatterplots, and 2x2 contingency tables.
- Relying only on bivariate correlations can mislead teams by ignoring lurking variables that distort the true relationship.

Diagnostic Level	Analytical Tools	Business Decision Action
Univariate	Histograms, Boxplots, Summary statistics	Identify outliers, flag invalid transaction values
Bivariate	Scatterplots, Cross-tabs, correlation coefficients	Formulate initial customer segmentation hypotheses
Key Risk	Correlation does not equal causation	Avoid launching marketing campaigns based on simple correlation

Multivariate & Hypothesis Scope

- Multivariate diagnostics analyze three or more variables to isolate cohort behaviors and control for collinearity.
- Hypothesis-driven exploration structures the analysis around business assumptions, preventing aimless data queries.
- Executives use these deep diagnostics to build high-fidelity simulation models for scenario planning and risk assessment.

Diagnostic Level	Analytical Tools	Business Value Impact
Multivariate	3D plots, Facet grids, Multi-way tables, Logistic regression	Isolate specific high-value customer cohorts
Hypothesis-Driven	MECE Issue trees, Value-driver trees, A/B test validation	Quantify EBITDA impact of changing operational levers
Execution Value	High-fidelity strategic planning	Target resources on high-impact operational improvements

Strategic Root Cause Example

- Problem: A retail firm faces a 4% drop in customer retention; bivariate analysis shows no correlation with price increases.
- Multivariate Audit: Cross-referencing retention with delivery delays and customer support wait times reveals a critical threshold.
- Business Resolution: Retaining customers requires prioritizing shipping carrier performance rather than cutting product prices.

Diagnostic Phase	Exploratory Level	Operational Discovery
Phase 1: Basic Audit	Level 1: Univariate	Overall customer churn rate has risen from 12% to 16%
Phase 2: Correlation Scan	Level 2: Bivariate	No significant correlation between customer churn and product price
Phase 3: Deep Cohort Audit	Level 3: Multivariate	Churn spikes to 48% only when delivery delay > 5 days AND wait time > 10m
Phase 4: Value Translation	Level 4: Hypothesis-Driven	Reducing shipping delays by 2 days saves $1.2M in annual churn revenue